


## Object Localization:



### output prediction vector:

[  $P_c$ ,  $b_x$ ,  $b_y$ ,  $b_h$ ,  $b_w$ ,  $c_1$ ,  $c_2$ ,  $c_3$  ]

#### ◆ Meaning of Each Component

Symbol	What it represents	Type
$P_c$	Objectness score — probability that an object exists in this cell	Scalar
$b_x$	Predicted bounding box <b>center x</b> relative to the cell	Scalar
$b_y$	Predicted bounding box <b>center y</b> relative to the cell	Scalar
$b_h$	Predicted bounding box <b>height</b>	Scalar
$b_w$	Predicted bounding box <b>width</b>	Scalar
$c_1$ , $c_2$ , $c_3$	Class probabilities (e.g., 3 classes)	Scalar s

So this vector has **1 + 4 + 3 = 8 values**.

### If $P_c = 1$

$P_c = 1$   
 $b_x = 0.4$   
 $b_y = 0.6$   
 $b_w = 0.3$   
 $b_h = 0.5$   
 $[c_1, c_2, c_3] = [0.1, 0.7, 0.2] \rightarrow$  Class 2 highest

### If $P_c = 0$

$P_c = 0$   
 $b_x, b_y, b_w, b_h =$  don't matter  
 $[c_1, c_2, c_3] =$  don't matter

### Square loss according to YOLOv1:


$$\text{Loss} = \lambda_{\text{coord}} * \sum[(b_x - \hat{b}_x)^2 + (b_y - \hat{b}_y)^2 + (b_w - \hat{b}_w)^2 + (b_h - \hat{b}_h)^2] \\ + (P_c - \hat{P}_c)^2 \\ + \sum_{\text{classes}}(c_i - \hat{c}_i)^2$$

Where:

Symbol	Meaning
$\lambda_{\text{coord}}$	Extra weight for box accuracy
$b_x, b_y, b_w, b_h$	Predicted box
$P_c$	Predicted object presence
$c_i$	Predicted class probabilities
$\hat{\cdot}$	Ground truth

Condition	What Loss Applies?	Why
$\hat{P}_c = 1$ (object exists)	Box loss + Class loss + $P_c$ loss	Model must learn full prediction
$\hat{P}_c = 0$ (no object)	Only $P_c$ loss	Model should predict background correctly

## Landmark Detection:



Landmark detection means identifying **specific key points** on an object (not just a bounding box).

Examples:

- Face: eyes, nose, mouth corners, chin
- Hand: finger joints
- Body: hips, knees, shoulders (pose estimation)
- Cars: wheel centers, headlights

## Output

```
[[left_eye_x, left_eye_y],  
 [right_eye_x, right_eye_y],  
 [nose_x, nose_y],  
 [mouth_left_x, mouth_left_y],  
 [mouth_right_x, mouth_right_y]]
```

## Convolutional Implementation Using Sliding Window for Object Detection

Classical object detection used a sliding window approach, where a fixed-size window moves across the image and each cropped region is sent to a classifier. This method is slow because the classifier repeatedly processes many overlapping regions.

Convolutional Neural Networks (CNNs) solve this inefficiency by converting the sliding-window operation into a single convolutional pass. In CNNs, the same set of filters slides over the image, meaning each convolution operation naturally evaluates local image regions—exactly like sliding windows. Fully connected layers of a classifier can be replaced by  $1 \times 1$  convolution layers, allowing the network to take the full image as input and output a heatmap where each cell represents the classifier's decision for each window location.

Thus, instead of cropping thousands of windows, the CNN performs all evaluations at once. This makes convolutional sliding-window detection fast, efficient, and the foundation for modern detectors such as Faster R-CNN, SSD, and YOLO.

## Intersection over Union (IoU)

Intersection over Union (IoU) is a metric used in object detection to measure how well a predicted bounding box matches the ground-truth box. It is defined as the ratio between the area of overlap and the area of union of the two boxes:

$$\text{Intersection}(\text{Predicted\_Box}, \text{Ground\_Truth\_Box})$$

$$\text{IoU} = \frac{\text{Intersection}(\text{Predicted\_Box}, \text{Ground\_Truth\_Box})}{\text{Area}(\text{Predicted\_Box}) + \text{Area}(\text{Ground\_Truth\_Box}) - \text{Intersection}}$$

The value ranges from 0 to 1, where:

- 1 means a perfect match between predicted and ground-truth boxes
- 0 means no overlap

## **Non-max Suppression:**



## **NMS Algorithm (Step-by-Step):**

1. Start with all predicted boxes and their confidence scores.
2. Select the box with the highest confidence score.
3. Remove all other boxes that have IoU greater than a chosen threshold (e.g., 0.5) with this box.
4. Repeat until no boxes remain.

## **Anchor Boxes**

Anchor boxes are predefined bounding boxes of different shapes and sizes placed at each location of a feature map in object detection models like YOLO, SSD, and Faster R-CNN.

They help the model detect objects of varying sizes and aspect ratios. For each anchor box, the model predicts:

- Adjustments to fit the object (position and size)
- Class probabilities

By using multiple anchor boxes per location, the detector can handle multiple objects in the same region and improve detection accuracy.

**Key point:** Anchor boxes act as templates that are adjusted to match the actual objects in the image.

**Keep in mind that a grid size of 13 \* 13 means dividing the image into 13 rows by 13 columns, and each grid cell is `image_shape / grid_size = pixels`. And modern object detection algorithms like YOLO don't use sliding window concept, a classical method.**

## YOLO Algo

1. Input image divided into a grid ( $S \times S$ )
2. Each grid cell predicts:
  - Bounding box → center ( $b_x, b_y$ ), width  $b_w$ , height  $b_h$
  - Objectness score  $P_c$  → how likely an object is present
  - Class probabilities for  $C$  classes
3. Confidence Score:

$$\text{Confidence} = P_c \times \text{IoU}(\text{pred}, \text{truth})$$

4. During inference:
  - Filter low-confidence boxes
  - Apply Non-Max Suppression (NMS) to remove duplicates

### ◆ Output Format per Box

$(P_c, b_x, b_y, b_w, b_h, p_1, p_2, \dots, p_C)$

Where:

- $P_c$  = Objectness/confidence score
- $b_x, b_y$  = Bounding box center coordinates
- $b_w, b_h$  = Bounding box width and height
- $p_1, p_2, \dots, p_C$  = Class probabilities for  $C$  classes

## ◆ Loss Function (Multi-Task)

The YOLO loss function consists of three main components:

- ✓ Localization loss (box center & size)
- ✓ Confidence loss
- ✓ Classification loss

Combined as a weighted sum to balance object vs. background detection:

$$\begin{aligned} \text{Loss} = & \lambda_{\text{coord}} * \sum [(b_x - \hat{b}_x)^2 + (b_y - \hat{b}_y)^2 + (b_w - \hat{b}_w)^2 + (b_h - \hat{b}_h)^2] \\ & + (P_c - \hat{P}_c)^2 \\ & + \sum_{\text{classes}} (c_i - \hat{c}_i)^2 \end{aligned}$$

---

## ◆ Anchors & Multi-Scale Detection (Modern YOLO)

Anchor boxes help detect objects of various shapes and sizes.

YOLOv3/v4/v5/v7/v8 use 3 detection scales:

Grid Size	Object Size	Use Case
13×13	Large objects	Detecting big objects in the scene
26×26	Medium objects	Detecting medium-sized objects
52×52	Small objects	Detecting fine-grained, small objects

